

# Introduction to Machine Learning

## Linear Regression

Varun Chandola

Computer Science & Engineering  
State University of New York at Buffalo  
Buffalo, NY, USA  
chandola@buffalo.edu



University at Buffalo  
Department of Computer Science  
and Engineering  
School of Engineering and Applied Sciences



## Basics

## Linear Regression

- Problem Formulation
- Learning Parameters
- Machine Learning as Optimization
- Convex Optimization
- Matrix Calculus Basics
- Gradient Descent
- Issues with Gradient Descent
- Stochastic Gradient Descent

- ▶ Data - scalar ( $x$ ), vector ( $\mathbf{x}$ ), Matrix ( $\mathbf{X}$ )

## Scalars

- ▶ Numeric ( $x \in \mathbb{R}$ )
- ▶ Categorical (e.g.,  $x \in \{0, 1\}$ )
- ▶ Constants will be denoted as  $D$ ,  $M$ , etc.

## Vector

- ▶ Length of a vector  $\mathbf{x} \in \mathbb{R}^D$
- ▶ Vector *dot* product ( $\mathbf{x} \cdot \mathbf{y}$ )
- ▶ Norm of a vector ( $|\mathbf{x}|$ ,  $\|\mathbf{x}\|$ ,  $\|\mathbf{x}\|_p$ )

## Matrix

- ▶ Size of a matrix ( $\mathbf{X} \in \mathbb{R}^{M \times N}$ )
- ▶ Transpose of a matrix ( $\mathbf{X}^\top$ )
- ▶ Matrix product ( $\mathbf{XY}$ )

- ▶ A vector is a special matrix with only one column

$$\mathbf{x} \cdot \mathbf{y} \equiv \mathbf{x}^\top \mathbf{y}$$

# Linear Regression

- ▶ There is one scalar **target** variable  $y$
- ▶ There is one vector **input** variable  $x$
- ▶ Inductive bias:

$$y = \mathbf{w}^\top \mathbf{x}$$

## Linear Regression Learning Task

Learn  $\mathbf{w}$  given training examples,  $\langle \mathbf{X}, \mathbf{y} \rangle$ .

# Geometric Interpretation

- ▶ Fitting a straight line to  $d$  dimensional data

$$y = \mathbf{w}^\top \mathbf{x}$$

$$y = \mathbf{w}^\top \mathbf{x} = w_1x_1 + w_2x_2 + \dots + w_dx_d$$

- ▶ Will pass through origin
- ▶ Add intercept

$$y = w_0 + w_1x_1 + w_2x_2 + \dots + w_dx_d$$

- ▶ Equivalent to adding another column in  $\mathbf{X}$  of 1s.

# Incorporating Bias/Intercept

## Explicit Bias

$$\mathbf{x} \equiv \{x_1, x_2, \dots, x_d\}$$

$$\mathbf{w} \equiv \{w_1, w_2, \dots, w_d\}$$

$$y = w_0 + \mathbf{w}^\top \mathbf{x}$$

## Implicit Bias

$$\mathbf{x} \equiv \{1, x_1, x_2, \dots, x_d\}$$

$$\mathbf{w} \equiv \{w_0, w_1, w_2, \dots, w_d\}$$

$$y = \mathbf{w}^\top \mathbf{x}$$

# Learning Parameters - Least Squares Approach

- ▶ Minimize *squared loss*

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y_i - \mathbf{w}^\top \mathbf{x}_i)^2$$

- ▶ or,

$$J(\mathbf{w}) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\mathbf{w})^\top (\mathbf{y} - \mathbf{X}\mathbf{w})$$

- ▶ Make prediction ( $\mathbf{w}^\top \mathbf{x}_i$ ) as close to the target ( $y_i$ ) as possible
- ▶ Least squares estimate

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$$

- ▶ *We will derive this expression in class.*

# Machine Learning as Optimization Problem<sup>1</sup>

- ▶ Learning is optimization
- ▶ Faster optimization methods for faster learning
- ▶ Let  $\mathbf{w} \in \mathbb{R}^d$  and  $S \subset \mathbb{R}^d$  and  $f_0(\mathbf{w}), f_1(\mathbf{w}), \dots, f_m(\mathbf{w})$  be real-valued functions.
- ▶ Standard optimization formulation is:

$$\begin{array}{ll}\underset{\mathbf{w}}{\text{minimize}} & f_0(\mathbf{w}) \\ \text{subject to} & f_i(\mathbf{w}) \leq 0, \quad i = 1, \dots, m.\end{array}$$

---

<sup>1</sup>Adapted from [http://ttic.uchicago.edu/~gregory/courses/ml2012/lectures/tutorial\\_optimization.pdf](http://ttic.uchicago.edu/~gregory/courses/ml2012/lectures/tutorial_optimization.pdf). Also see, <http://www.stanford.edu/~boyd/cvxbook/> and [http://scipy-lectures.github.io/advanced/mathematical\\_optimization/](http://scipy-lectures.github.io/advanced/mathematical_optimization/).




# Solving Optimization Problems


- ▶ Methods for **general optimization problems**
  - ▶ Simulated annealing, genetic algorithms
- ▶ Exploiting *structure* in the optimization problem
  - ▶ **Convexity**, Lipschitz continuity, smoothness

# Convexity

## Convex Sets



## Convex Functions



- Optimality Criterion

$$\begin{array}{ll}\underset{\mathbf{w}}{\text{minimize}} & f_0(\mathbf{w}) \\ \text{subject to} & f_i(\mathbf{w}) \leq 0, \quad i = 1, \dots, m.\end{array}$$

where all  $f_i(\mathbf{w})$  are **convex functions**.

- $\mathbf{w}_0$  is feasible if  $\mathbf{w}_0 \in \text{Dom } f_0$  and all constraints are satisfied
- A feasible  $\mathbf{w}^*$  is optimal if  $f_0(\mathbf{w}^*) \leq f_0(\mathbf{w})$  for all  $\mathbf{w}$  satisfying the constraints

# Matrix Calculus Basics

$$\frac{\partial \mathbf{a}^\top \mathbf{b}}{\partial \mathbf{a}} = \frac{\partial \mathbf{b}^\top \mathbf{a}}{\partial \mathbf{a}} = \mathbf{b}$$


$$\frac{\partial \mathbf{a}^\top \mathbf{M} \mathbf{a}}{\partial \mathbf{a}} = 2\mathbf{M} \mathbf{a}$$

where  $\mathbf{M}$  is a symmetric matrix.

# Gradient of a Function

- Denotes the direction of steepest ascent

$$\nabla J(\mathbf{w}) = \begin{bmatrix} \frac{\partial J}{\partial w_0} \\ \frac{\partial J}{\partial w_1} \\ \vdots \\ \frac{\partial J}{\partial w_d} \end{bmatrix}$$



# Finding Extremes of a Single Variable Function

1. Set derivative to 0

$$\nabla J(\mathbf{w}) = 0$$

2. Check second derivative for minima or maxima or saddle point

# Finding Extremes of a Multiple Variable Function - Gradient Descent

1. Start from any point in variable space
2. Move along the direction of the steepest descent (or ascent)
  - ▶ By how much?
  - ▶ A learning rate ( $\eta$ )
  - ▶ What is the direction of steepest descent?
    - ▶ Gradient of  $J$  at  $\mathbf{w}$

## Training Rule for Gradient Descent

$$\mathbf{w} = \mathbf{w} - \eta \nabla J(\mathbf{w})$$

For each weight component:

$$w_j = w_j - \eta \frac{\partial J}{\partial w_j}$$

# Convergence Guaranteed?

- ▶ Error surface contains only one global minimum
- ▶ Algorithm *will* converge
  - ▶ Examples need not be linearly separable
- ▶  $\eta$  should be *small enough*
- ▶ Impact of too large  $\eta$ ?
- ▶ Too small  $\eta$ ?



# Issues with Gradient Descent

- ▶ Slow convergence
- ▶ Stuck in local minima

# Stochastic Gradient Descent [1]

- ▶ **Update weights after every training example.**
- ▶ For sufficiently small  $\eta$ , closely approximates Gradient Descent.

<b>Gradient Descent</b>	<b>Stochastic Gradient Descent</b>
Weights updated after summing error over all examples	Weights updated after examining each example
More computations per weight update step	Significantly lesser computations
Risk of local minima	Avoids local minima

# Gradient Descent Based Method

- ▶ Minimize the squared loss using *Gradient Descent*

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^N (y_i - \mathbf{w}^\top \mathbf{x}_i)^2$$

- ▶ Why?

# References



Y. LeCun, B. Boser, J. S. Denker, D. Henderson, R. E. Howard, W. Hubbard, and L. D. Jackel.

Backpropagation applied to handwritten zip code recognition.  
*Neural Comput.*, 1(4):541–551, Dec. 1989.