

# Introduction to Machine Learning

## Neural Networks

Varun Chandola

Computer Science & Engineering  
State University of New York at Buffalo  
Buffalo, NY, USA  
chandola@buffalo.edu



University at Buffalo  
Department of Computer Science  
and Engineering  
School of Engineering and Applied Sciences



Extending Linear Models

Multi Layered Perceptrons

Generalizing to Multiple Labels

Properties of Sigmoid Function

Feed Forward Neural Networks

Backpropagation

Derivation of the Backpropagation Rules

Final Algorithm

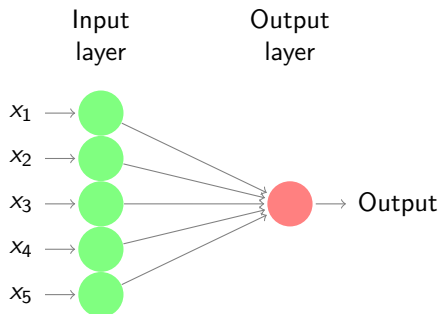
Wrapping up Neural Networks

Bias Variance Tradeoff


# Extending Linear Models

## ► Questions?

- How to learn non-linear surfaces?
- How to generalize to multiple outputs, numeric output?



# Anatomy of a Unit ( $r$ )




# Generalizing to Multiple Labels

- ▶ Distinguishing between multiple categories
- ▶ *Solution:* Add another layer - **Multi Layer Neural Networks**

# Generalizing to Multiple Labels

- ▶ Distinguishing between multiple categories
- ▶ *Solution:* Add another layer - **Multi Layer Neural Networks**



# What Activation Unit to Use?


- ▶ Linear Unit
- ▶ Perceptron Unit

# What Activation Unit to Use?

- ▶ Linear Unit
- ▶ Perceptron Unit
- ▶ Sigmoid Unit
  - ▶ Smooth, differentiable activation function

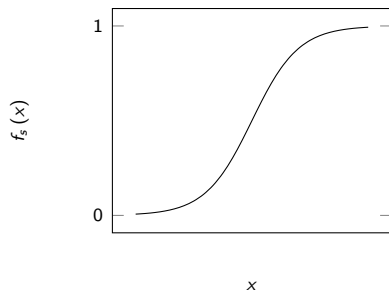
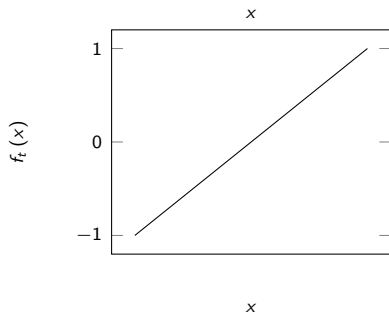
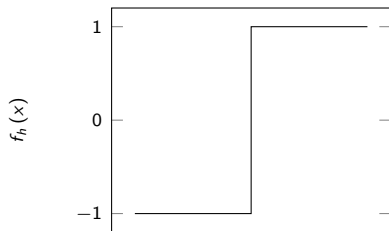
$$\sigma(net) = \frac{1}{1 + e^{-net}}$$

- ▶ Non-linear output






# Properties of Sigmoid Function



# Feed Forward Neural Networks - Architecture



# Feed Forward Neural Networks

- ▶  $D$  input nodes (excluding bias)
  - ▶  $M$  hidden nodes (excluding bias)
  - ▶  $K$  output nodes
- ▶ At hidden nodes:  $\mathbf{w}_j, 1 \leq j \leq M, \mathbf{w}_j \in \mathbb{R}^{D+1}$
  - ▶ At output nodes:  $\mathbf{w}_l, 1 \leq l \leq K, \mathbf{w}_l \in \mathbb{R}^{M+1}$

# Learning Weights of the Multi-layer Network

- ▶ Assume that the network structure is predetermined (number of hidden nodes and interconnections)
- ▶ Objective function for  $N$  training examples:

$$J = \sum_{i=1}^N J_i = \frac{1}{2} \sum_{i=1}^N \sum_{l=1}^K (y_{il} - o_{il})^2$$

- ▶  $y_{il}$  - Target value associated with  $l^{th}$  class for input ( $\mathbf{x}_i$ )
- ▶  $y_{il} = 1$  when  $k$  is true class for  $\mathbf{x}_i$ , and 0 otherwise
- ▶  $o_{il}$  - Predicted output value at  $l^{th}$  output node for  $\mathbf{x}_i$

What are we learning?

Weight vectors for all output and hidden nodes that minimize  $J$

# The Backpropagation Algorithm

1. Initialize all weights to *small values*
2. For each training example,  $\langle \mathbf{x}, \mathbf{y} \rangle$ :
  - 2.1 **Propagate input forward** through the network
  - 2.2 **Propagate errors backward** through the network

# Backpropagation Algorithm - Continued

## Gradient Descent


- ▶ Move in the opposite direction of the **gradient** of the objective function
- ▶  $-\eta \nabla J$

$$\nabla J = \sum_{i=1}^N \nabla J_i$$

- ▶ What is the gradient computed with respect to?
  - ▶ Weight vectors -  $M$  at hidden nodes and  $K$  at output nodes
  - ▶  $\mathbf{w}_j$  ( $j = 1 \dots M$ )
  - ▶  $\mathbf{w}_l$  ( $l = 1 \dots K$ )
- ▶  $\mathbf{w}_j \leftarrow \mathbf{w}_j - \eta \frac{\partial J}{\partial \mathbf{w}_j} = \mathbf{w}_j - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial \mathbf{w}_j}$
- ▶  $\mathbf{w}_l \leftarrow \mathbf{w}_l - \eta \frac{\partial J}{\partial \mathbf{w}_l} = \mathbf{w}_l - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial \mathbf{w}_l}$

$$\nabla J_i = \begin{bmatrix} \frac{\partial J_i}{\partial \mathbf{w}_1} \\ \frac{\partial J_i}{\partial \mathbf{w}_2} \\ \vdots \\ \frac{\partial J_i}{\partial \mathbf{w}_{m+k}} \end{bmatrix}$$

# Anatomy of a Sigmoid Unit ( $r$ )



$$\frac{\partial J_i}{\partial \mathbf{w}_r} = \begin{bmatrix} \frac{\partial J_i}{\partial w_{r1}} \\ \frac{\partial J_i}{\partial w_{r2}} \\ \vdots \end{bmatrix}$$

- ▶ Need to compute  $\frac{\partial J_i}{\partial w_{rq}}$
- ▶ Update rule for the  $q^{th}$  entry in the  $r^{th}$  weight vector:

$$w_{rq} \leftarrow w_{rq} - \eta \frac{\partial J}{\partial w_{rq}} = w_{rq} - \eta \sum_{i=1}^N \frac{\partial J_i}{\partial w_{rq}}$$

# Derivation of the Backpropagation Rules

Assume that we only one training example, i.e.,  $i = 1$ ,  $J = J_i$ . Dropping the subscript  $i$  from here onwards.

- ▶ Consider any weight  $w_{rq}$
- ▶ Let  $u_{rq}$  be the  $q^{th}$  element of the input vector coming in to the  $r^{th}$  unit.

## Observation 1

Weight  $w_{rq}$  is connected to  $J$  through  $net_r = \sum_q w_{rq} u_{rq}$ .

$$\frac{\partial J}{\partial w_{rq}} = \frac{\partial J}{\partial net_r} \frac{\partial net_r}{\partial w_{rq}} = \frac{\partial J}{\partial net_r} u_{rq}$$



# Analyzing Output Nodes

## Observation 2

$net_I$  for an **output node** is connected to  $J$  only through the output value of the node (or  $o_I$ )

$$\frac{\partial J}{\partial net_I} = \frac{\partial J}{\partial o_I} \frac{\partial o_I}{\partial net_I}$$

# Analyzing Output Nodes

## Observation 2

$net_I$  for an **output node** is connected to  $J$  only through the output value of the node (or  $o_I$ )

$$\frac{\partial J}{\partial net_I} = \frac{\partial J}{\partial o_I} \frac{\partial o_I}{\partial net_I}$$

## Update Rule for Output Units

$$w_{Ij} \leftarrow w_{Ij} + \eta \delta_I u_{Ij}$$

where  $\delta_I = (y_I - o_I)o_I(1 - o_I)$ .

► *Question:* What is  $u_{Ij}$  for the  $I^{th}$  output node?

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

# Analyzing Hidden Nodes

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

## Update Rule for Hidden Units

$$w_{jp} \leftarrow w_{jp} + \eta \delta_j u_{jp}$$

# Analyzing Hidden Nodes

## Observation 3

$net_j$  for a **hidden node** is connected to  $J$  through all output nodes

$$\frac{\partial J}{\partial net_j} = \sum_{l=1}^K \frac{\partial J}{\partial net_l} \frac{\partial net_l}{\partial net_j}$$

## Update Rule for Hidden Units

$$w_{jp} \leftarrow w_{jp} + \eta \delta_j u_{jp}$$

$$\delta_j = o_j(1 - o_j) \sum_{l=1}^K \delta_l w_{lj}$$

$$\delta_l = (y_l - o_l) o_l(1 - o_l)$$

- *Question:* What is  $u_{jp}$  for the  $j^{th}$  hidden node?

- ▶ While not converged:
  - ▶ *Move forward* to compute outputs at hidden and output nodes
  - ▶ *Move backward* to propagate errors back
    - ▶ Compute  $\delta$  errors at output nodes ( $\delta_l$ )
    - ▶ Compute  $\delta$  errors at hidden nodes ( $\delta_j$ )
  - ▶ Update all weights according to weight update equations

# Conclusions about Neural Networks

- ▶ Error function contains many local minima
- ▶ No guarantee of convergence
  - ▶ Not a “big” issue in practical deployments
- ▶ Improving backpropagation

# Conclusions about Neural Networks

- ▶ Error function contains many local minima
- ▶ No guarantee of convergence
  - ▶ Not a “big” issue in practical deployments
- ▶ Improving backpropagation
  - ▶ Adding momentum
  - ▶ Using stochastic gradient descent
  - ▶ Train multiple times using different initializations



# Bias Variance Tradeoff

- ▶ Neural networks are *universal function approximators*
  - ▶ By making the model more complex (increasing number of hidden layers or  $m$ ) one can lower the error
- ▶ Is the model with least training error the best model?

# Bias Variance Tradeoff

- ▶ Neural networks are *universal function approximators*
  - ▶ By making the model more complex (increasing number of hidden layers or  $m$ ) one can lower the error
- ▶ Is the model with least training error the best model?
  - ▶ The simple answer is **no!**
  - ▶ Risk of overfitting (chasing the data)
  - ▶ Overfitting  $\Leftarrow$  **High generalization error**

## High Variance - Low Bias

- ▶ “Chases the data”
- ▶ Very low training error
- ▶ Poor performance on unseen data

## Low Variance - High Bias

- ▶ Less sensitive to training data
- ▶ Higher training error
- ▶ Better performance on unseen data

# Getting the Right Balance

- ▶ General rule of thumb – If two models are giving similar training error, choose the **simpler** model
- ▶ What is simple for a neural network?
- ▶ Low weights in the weight matrices?
  - ▶ Why?

# Introducing Bias in Neural Network Training

- ▶ Penalize solutions in which the weights are high
- ▶ Can be done by introducing a penalty term in the objective function
  - ▶ **Regularization**

## Regularization for Backpropagation

$$\tilde{J} = J + \frac{\lambda}{2n} \left( \sum_{j=1}^M \sum_{i=1}^{D+1} (w_{ji}^{(1)})^2 + \sum_{l=1}^K \sum_{j=1}^{M+1} (w_{lj}^{(2)})^2 \right)$$

# Other Extensions?

- ▶ Use a different loss function (why)?
  - ▶ Quadratic (Squared), Cross-entropy, Exponential, KL Divergence, etc.
- ▶ Use a different activation function (why)?
  - ▶ Sigmoid

$$f(z) = \frac{1}{1 + \exp(-z)}$$

- ▶ Tanh

$$f(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- ▶ Rectified Linear Unit (ReLU)

$$f(z) = \max(0, z)$$

# References